
Satellite Computer Vision mit Keras und TensorFlow

Best Practices und Beispiele aus der Forschung

Dr.-Ing. Jens Leitloff
Felix M. Riese
Mannheim, May 2019



Why

is satellite computer vision relevant?

How

does computer vision help remote sensing?

What

do you need to start?

Outline

Why

is satellite computer vision relevant?

How

does computer vision help remote sensing?

What

do you need to start?

→ Theory

~ 10 minutes

GitHub



→ Theory

→ Code examples in slides

→ Research results

~ 25 minutes

→ Best practices

→ Links to resources

~ 5 minutes

jensleitloff/CNN-Sentinel

Why

is satellite computer vision relevant?

How

does computer vision help remote sensing?

What

do you need to start?

1. Why is this talk relevant (for me)?
2. What is satellite CV?
3. How is it applied in today's life?

Why is this talk relevant (for me)?



Engineers & scientists

State-of-the-art methods
Smart data interpretation
Publishing methods & code



Similar challenges

Multi-modal data fusion
Weak & sparse labels
Curse of dimensionality




Relevant data

2D, 3D, time series, GIS
A lot of free datasets
Big data & small data

Remote sensing & satellites

Active


- Microwave (SAR)
- Laser (LIDAR)
- Acoustic (SONAR)



Passive

- RGB
- Multispectral

→ Focus of this talk




Applications of satellite remote sensing


- **Meteorology** → weather forecasts, climate research, monitoring of glaciers and sea ice
- **Natural resources** → agriculture, forestry, environmental monitoring, coastal monitoring
- **Geosciences** → mapping, road networks, hydrology
- **Land administration** → land use / cover, building detection
- **Disaster response** → monitoring of natural hazards like wildfires, floods, earthquakes, volcanoes



Example: Land cover



Example: Land cover & passive satellite





The Copernicus programme (by ESA)

Launch date	Mission	Objective
2014	Sentinel-1	Radar for land and ocean
2015	Sentinel-2	Vegetation and soil monitoring
2016	Sentinel-3	Marine observation
→ 2023	Sentinel-4	Air quality monitoring
→ 2021	Sentinel-5	Air quality monitoring
→ 2020	Sentinel-6	Atmospheric measurements

5 days / world

It takes the Sentinel-2 satellites five days to cover the earth's surface.



3.2 TB / day

The Sentinel-2 satellites alone produce images in the size of 3.2 Terabytes per day, 365 days a year.



0€


The data of the Sentinel program is free for everyone.
No matter if commercial or not.
No matter if European or not.






Computer Vision

To bridge the gap between pixels
and “meaning”.



Computer Vision in Satellite Remote Sensing

To bridge the gap between pixels and “meaning”.



Why

is satellite computer vision relevant?

How

does computer vision help remote sensing?

What

do you need to start?

1. Learning 2D patterns with CNN
2. Learning 1D & 3D patterns
3. Learning sequences

Learning 2D patterns

The EuroSAT dataset

Links:

- Introduced in [arXiv:1709.00029](https://arxiv.org/abs/1709.00029)
- Download here: <http://madm.dfki.de/downloads>

Properties:

- Sentinel-2 satellite images
 - Raw dataset: 13 different spectral bands (1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 8a)
 - RGB dataset: RGB colors
- 27 000 labeled images
- 10 classes (→ next slide)




Image source: <https://www.cls.fr/en/sentinel-2-environmental-studies/>



(a) Industrial



(b) Residential



(c) Annual crop



(d) Permanent crop



(e) River



(f) Sea & lake




(g) Herbaceous vegetation



(h) Highway



(i) Pasture




(j) Forest

one class per image

Image source: [arXiv:1709.00029](https://arxiv.org/abs/1709.00029)

Segmentation training

one class per pixel



Learning from scratch

1. Network structure with **randomly initialized weights**
2. Training of **complete** network on training dataset

- slow convergence
- lots of training data necessary
- easily implemented for **every kind** of input data

Transfer learning

Learning from scratch

1. Network structure with **randomly initialized weights**
2. Training of **complete** network on training dataset

- slow convergence
- lots of training data necessary
- easily implemented for **every kind** of input data


Transfer learning

1. Network structure with **pretrained weights**
2. Training of **selected layers** on the training dataset (= fine tuning)


- much **faster** convergence
- **less** training data necessary
- mostly implemented for **RGB**
- modifications for MS necessary

→ Implement for RGB & multispectral

VGG16 architecture



DenseNet architecture





Steps to do

Training from scratch

1. Initialize network model without top layers
2. define new top layers
3. define data augmentation
4. define callbacks
5. -
6. -
7. -
8. fit model

Transfer-learning

1. Pretrained network model without top layers
2. define new top layers
3. define data augmentation
4. define callbacks
5. set base layers non trainable
6. fit model (train new top layers)
7. set (some) base layers trainable
8. fit model (fine-tune base and top layers)

Learning from scratch

```
base_model = VGG(include_top=False,  
                  weights=None,  
                  input_shape=(64, 64, 13))
```

```
# add a global spatial average pooling layer  
top_model = base_model.output  
top_model = GlobalAveragePooling2D()(top_model)  
# let's add a fully-connected layer  
top_model = Dense(2048, activation='relu')(top_model)  
top_model = Dense(2048, activation='relu')(top_model)  
# and a logistic layer  
predictions = Dense(num_classes, activation='softmax')(top_model)  
  
# this is the model we will train  
model = Model(inputs=base_model.input, outputs=predictions)
```

Transfer learning

```
base_model = VGG(include_top=False,  
                  weights='imagenet',  
                  input_shape=(64, 64, 3))
```

```
# first: train only the top layers  
# i.e. freeze all convolutional layers  
for layer in base_model.layers:  
    layer.trainable = False
```

Nice tools in Keras I

- Image augmentation
- Use image patches in directories

```
# defining ImageDataGenerators
# ... initialization for training
train_datagen = ImageDataGenerator(fill_mode="reflect",
                                    rotation_range=45,
                                    horizontal_flip=True,
                                    vertical_flip=True)

train_datagen.flow_from_directory(path_to_train,
                                  target_size=(64, 64),
                                  batch_size=batch_size,
                                  class_mode='categorical')
```



Nice tools in Keras II

- Define some callbacks for training

```
checkpointer = ModelCheckpoint(...)  
earlystopper = EarlyStopping(...)  
tensorboard = TensorBoard(...)
```

- and train ...

```
model.fit_generator(train_generator,...  
                    callbacks=[checkpointer, earlystopper,  
                               tensorboard],  
                    validation_data=validation_generator)
```


Additional things in the repository

- in detail notebooks for
 - learning from scratch
 - transfer learning
 - customized data generators
- python scripts for
 - panchromatic data
 - multispectral data
- many additional information about
 - setup an environment
 - data
 - ...

GitHub




jensleitloff/CNN-Sentinel




Karlsruhe, August 2016

30




Annual Crop
Forest
Herbaceous Vegetation
Highway
Industrial
Pasture
Permanent Crop
Residential
River
SeaLake




Karlsruhe, August 2018

Ludwigshafen, August 2015



Results from project with LVerMGeo Rheinland-Pfalz



Results with image segmentation



Klassifikation 20cm

Learning 1D & 3D patterns

Learning sequences




Other convolution approaches

Convolution approaches:

- No convolution
- 1D convolution over spectral dimension
- 2D convolution over spatial dimension → this talk
- 3D convolution over spectral & spatial dimension
- 3D convolution over temporal & spatial dimension

Review paper: [arXiv:1710.03959](https://arxiv.org/abs/1710.03959)

Learning 1D patterns with CNNs



- Paper will be presented in June 2019
- Code available on GitHub
- Free soil texture dataset
- Hyperspectral → more features than multispectral
- 1D convolution over the spectral dimension

Felix M. Riese and Sina Keller, "Soil Texture Classification with 1D Convolutional Neural Networks based on Hyperspectral Data," 2019, Accepted at Geospatial Week 2019 in Enschede (NL). Available: [arXiv:1901.04846](https://arxiv.org/abs/1901.04846)

Learning 3D patterns with CNNs

- 3D convolution over **spatial & spectral** dimension
 - spatial dimension → earth surface (cf. [Learning 2D patterns](#))
 - spectral dimension → making use of multi-/hyperspectral data (cf. [Learning 1D patterns with CNNs](#))
- 3D convolution over **spatial & temporal** dimension
 - temporal dimension → learn from temporal changes, e.g. crop

Learning sequences




- Temporal changes → learning sequences
- Recurrent neural networks, LSTM → Rußwurm, Körner ([arXiv:1802.02080](https://arxiv.org/abs/1802.02080))

One more thing: SUSI

SUervised Self-organizing maps in Python

- Self-organizing maps for unsupervised **and supervised** learning
- Code and usage in **scikit-learn style**
- Example applications
 - 2D visualization of high-dimensional data
 - Classification and regression on small datasets
- Code and paper freely available

see talk by Oliver
Zeigermann today



Source: [felixriese/susi](#)

Why

is satellite computer vision relevant?

How

does computer vision help remote sensing?

What

do you need to start?

1. Best practices from research
2. Starting point with satellite CV

Basics

Best practices

- Use existing frameworks like **Keras**
- Make use of **build-in functions** like
 - data augmentation generators
 - callbacks
- Try **established architectures** and try **pre-trained** networks to reduce training time
- **Big data beats data cleaning**
- Measure error on training **and test** subset
- Training **strategy**:
 1. Training of estimator on small dataset until training error is 0
 2. Implementation of regularization, e.g. dropout, batch-normalization, ...



Best practices

Further reading

Recommended resources:

- [CPVR 2014 Talk by Marc'Aurelio Ranzato](#)
- “Visualizing and Understanding Convolutional Networks” by M. D. Zeiler and R. Fergus, 2013, [arXiv:1311.2901](#)
- Coursera Deep Learning Specialization part 3: [Structuring Machine Learning Projects](#) by deeplearning.ai

Starting point with satellite CV

GitHub



jensleitloff/CNN-Sentinel

- Requirements (Versions ...)
- Environment setup (Conda ...)
- Code for
 - setting up the data
 - CNN transfer learning RGB & MS
 - CNN learning-from-scratch RGB & MS
- How to download Sentinel-2 data
- Links to more datasets



Jens Leitloff

✉ jens.leitloff@kit.edu
jensleitloff/CNN-Sentinel



Felix M. Riese

✉ felix.riese@kit.edu

Thank you for your attention!
